



Cooking Infrastructure isn't illegal

aka. How not to Break things Bad



Startup Safari 2018



Agenda

- Introductions
- Bad Infrastructure
- How Not To Cook
- Enter AWS
- CloudFormation
- Managing CloudFormation stacks - Furnace
- Live Coding
- Why Furnace



Introductions



Bad Infrastructure



Common Patterns

- Un-reproducible
- Fragile
- Difficult to maintain
- Has leprechauns (hand made changes)
- Too many "Why is this here? What does it do?" moments



How not to cook



Gordon Ramsay would be sad

- Using various recipes (have puppet and chef and ansible at the same time)
- Overcook (you infrastructure code is too complicated)
- Though, that's how pizza was made (have bread and flavour it with whatever you can find)



Enter AWS



Little History



Basement Servers

aka. the great sysadmin wars of 2001



<https://i.ytimg.com/vi/SVRT5TTr-Bk/maxresdefault.jpg>

You don't need these unless you really really REALLY REEEAALLLYYY want your data locally / you're a bank



Using AWS



Bad usage

- Having a single, monolith EC2 instance with
- No VPC
- No Security Groups
- No Configuration management to manage it

Might as well stick to Linode / Digital Ocean at that point.



Better usage

- Still, single EC2, but it's managed by a configuration management system
- Ansible, Chef, Puppet, Salt...
- But...
- No VPC
- No Security Groups
- No real single view of how many components there are




Good usage

- Security Groups and VPC applied
- Optionally: LoadBalancers, Autoscaling Groups
- Things are composed through some kind of configuration management
- Still... Things might be hand crafted / created. For example: Security groups, IAMs, VPCs or other AWS resources



Perfection...





CloudFormation

<https://aws.amazon.com/cloudformation/>



Idempotent

- What is CloudFormation?
- How can it help?
- Why would I use it?



Managing CloudFormation stacks - Furnace

<https://github.com/Skarlso/go-furnace>



Furnace



Configuring Furnace

```
main:
  stackname: FurnaceStack
  spinner: 1
aws:
  code_deploy_role: CodeDeployServiceRole
  region: us-east-1
  enable_plugin_system: false
  template_name: cloud_formation.template
  app_name: furnace_app
  code_deploy:
    # Only needed in case S3 is used for code deployment
    code_deploy_s3_bucket: furnace_code_bucket
    # The name of the zip file in case it's on a bucket
    code_deploy_s3_key: furnace_deploy_app
    # In case a Git Repository is used for the application, define these two settings
    git_account: Skarlso/furnace-codedeploy-app
    git_revision: b89451234...
```



Configuring a Stack

```
stacks/mystack.yaml
```



Create Stacks with Furnace

```
faws create
```



Deploy Applications with Furnace

```
faws push
```



Using Multiple Furnace files

```
- webstack
|-- webstack.yaml
|-- webstack.template
|
- databasestack
|-- database_stack.yaml
|-- database_stack.template
|
- .webstack.furnace
- .dbstack.furnace
```




Creating both stacks

```
faws create webstack  
faws create dbstack
```



Live coding

Commence some magic here



In case live coding fails...

Output of Create

```
[/] Waiting for state: CREATE_COMPLETE  
... Post-create plugin events...  
... Stack state is: CREATE_COMPLETE
```



Available maintenance commands

create

update

status

delete

push

delete-application



Why Furnace?



- * Single binary
- * No dependencies
- * Dead simple to use



**This isn't about
Furnace**



Thank you!

<https://github.com/Skarlso/go-furnace>